


# Apache Spark Installation & RDD

亦思科技 鄭紹志  
vito@is-land.com.tw

# About me

- 鄭紹志 Vito
- 任職亦思科技
- Hadoop/HBase/Spark 相關的研究開發工作
- 目前投入的領域主要為 Scala & Spark



歡迎加入 Spark@Hsinchu !

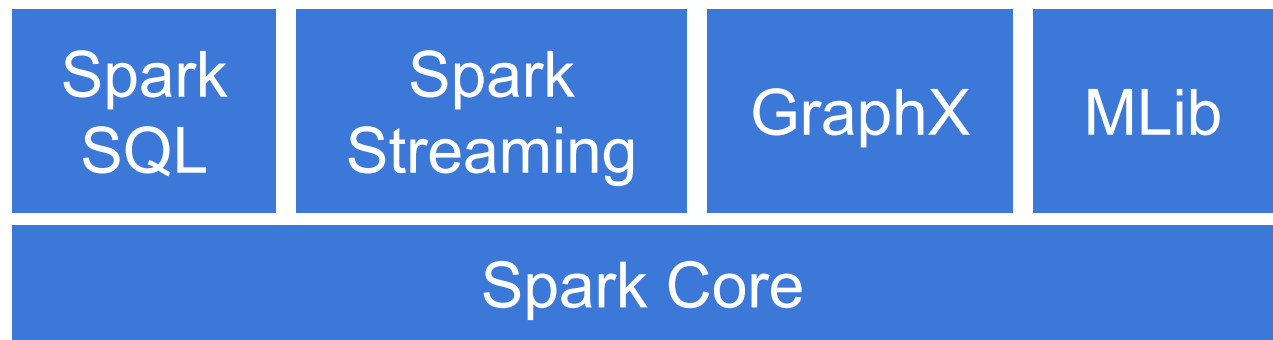
- Meetup -  
<https://www.meetup.com/Apache-Spark-Hsinchu/>
- Gitter 聊天室 -  
<https://gitter.im/hubertfc/SparkHsinchu>

# Agenda

- Spark overview
- Standalone cluster installation
- Spark Glossary
- RDD concept & overview
- RDD operations
- Shuffle & RDD Dependency

# Spark Overview

Apache Spark is a fast and general-purpose [\*cluster computing system\*](#). It provides high-level APIs in Java, Scala, Python and R, and an optimized engine that supports general execution graphs. It also supports a rich set of higher-level tools including [Spark SQL](#) for SQL and structured data processing, [MLlib](#) for machine learning, [GraphX](#) for graph processing, and [Spark Streaming](#).



# Cluster Mode

- Spark 支援三種叢集模式： standalone, mesos, yarn
- **Standalone** ✓
  - A simple cluster manager included with Spark
  - 最簡單的 Spark 叢集模式
  - 可以有單機版的 cluster - 方便開發測試
- **Apache Mesos**
  - A cluster manager. Enable resource sharing, improving cluster utilization.
  - <http://mesos.apache.org/>
- **Hadoop YARN**
  - A resource manager in Hadoop 2
  - <https://hadoop.apache.org/docs/r2.7.2/hadoop-yarn/hadoop-yarn-site/YARN.html>

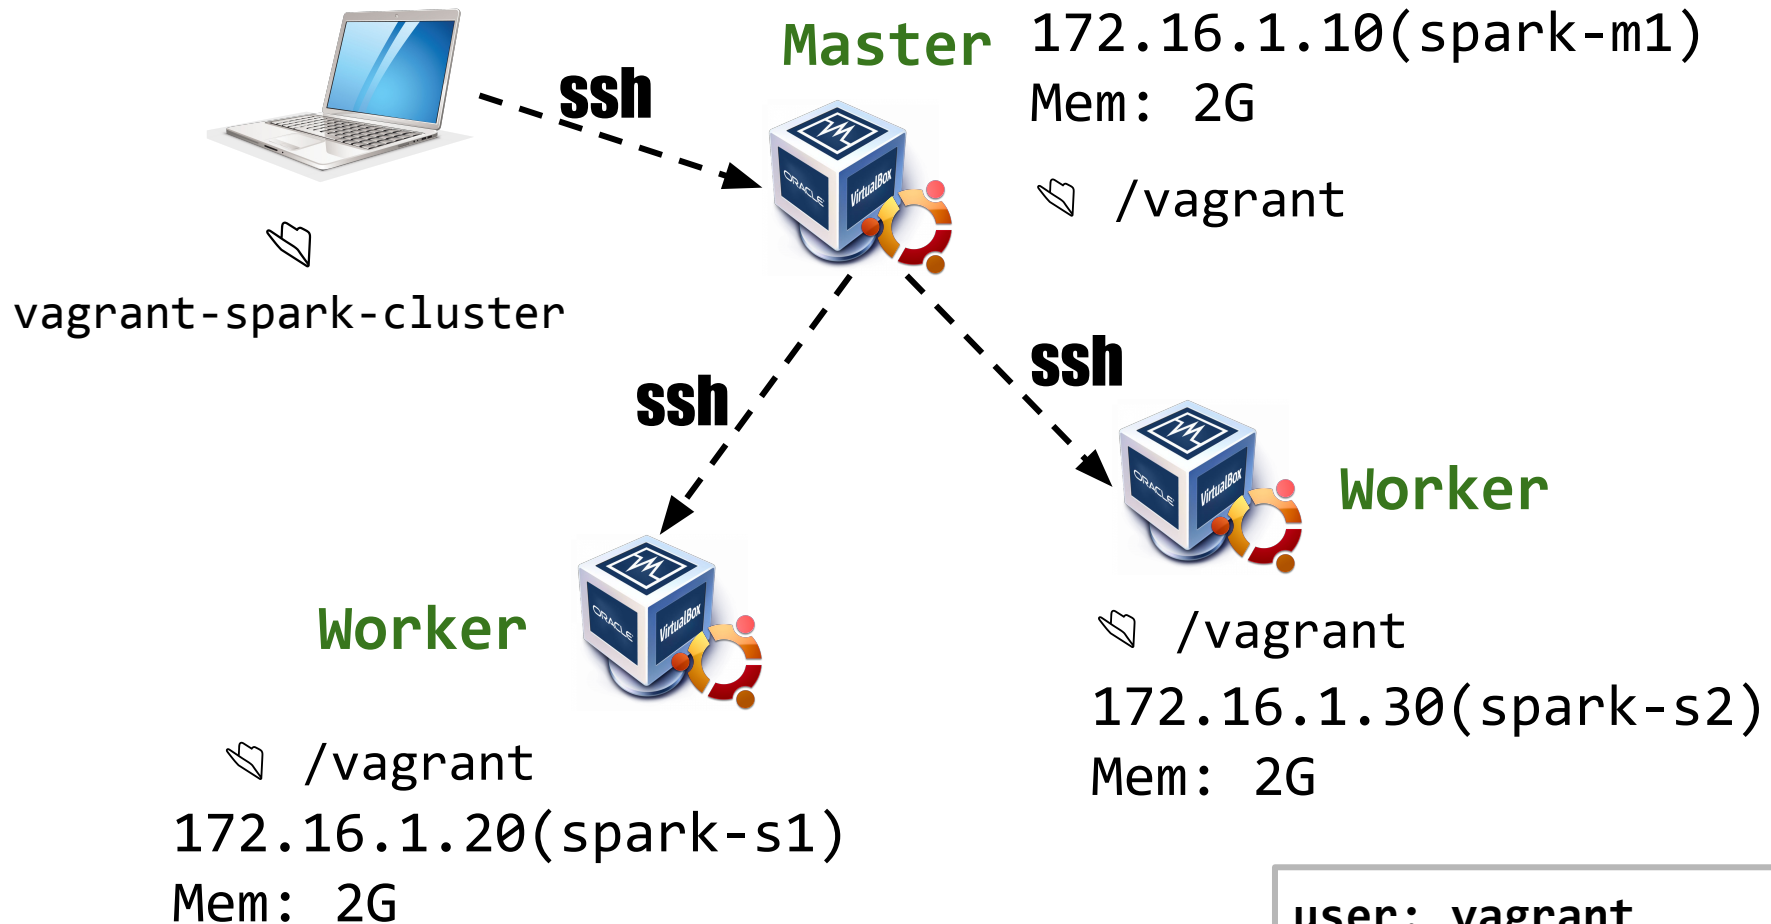
# Standalone cluster installation

# Spark standalone cluster

- 手工打造
- 自動化佈署 Spark: VirtualBox + Vagrant + Ansible
  - <https://github.com/is-land/vagrant-spark-cluster>
  - Cluster Information
    - OS: Ubuntu 14.04 LTS
    - JDK 1.8
    - Spark 1.6.2 Pre-Build for Hadoop 2.6
    - 1 Master node
    - 2 Slave node
    - Share folder: /vagrant
    - Master ssh login to Slave without password

<http://spark.apache.org/docs/1.6.2/spark-standalone.html>


# Apache Spark standalone cluster



user: vagrant  
password: vagrant



# Spark cluster architecture



# Spark Installation - Master(spark-m1)

- ssh 登入 **spark-m1**
- 安裝 JDK 1.8.x
  - 設定環境變數 JAVA\_HOME
  - 將 \$JAVA\_HOME/bin 加入 PATH 環境變數
- 修改 /etc/hosts, 加入以下資訊

```
172.16.1.10    spark-m1
172.16.1.20    spark-s1
172.16.1.30    spark-s2
```

- 產生 ssh key, 並複製到共享資料夾(/vagrant)內

```
vagrant@spark-m1:~$ ssh-keygen -t rsa -N "" -f "/home/vagrant/.ssh/id_rsa"
vagrant@spark-m1:~$ cp /home/vagrant/.ssh/id_rsa /vagrant/.ssh/id_rsa.master
vagrant@spark-m1:~$ cp /home/vagrant/.ssh/id_rsa.pub /vagrant/.ssh/id_rsa.master.pub
```

- 設定免密碼登入(**spark-m1** --> **spark-m1**)

```
vagrant@spark-m1:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
vagrant@spark-m1:~$ sudo su -
root@spark-m1:~$ cat /vagrant/.ssh/id_rsa.master.pub >> ~/.ssh/authorized_keys
```

# Spark Installation - Master(spark-m1)

- 下載 & 解壓縮 Spark 1.6.2(pre-build for Hadoop 2.6)

```
vagrant@spark-m1:~$ wget http://d3kbcqa49mib13.cloudfront.net/spark-1.6.2-bin-hadoop2.6.tgz
vagrant@spark-m1:~$ tar zxvf spark-1.6.2-bin-hadoop2.6.tgz
```

- 修改 ~/.profile 加入環境變數

```
export SPARK_HOME=/home/vagrant/spark-1.6.2-bin-hadoop2.6
```

- 登出 ssh, 再次登入 spark-m1. 並執行 Spark master server


```
vagrant@spark-m1:~$ cd spark-1.6.2-bin-hadoop2.6
vagrant@spark-m1:~/spark-1.6.2-bin-hadoop2.6$ sbin/start-master.sh
```

- 檢查是否正常啟動(process & log)

```
vagrant@spark-m1:~/spark-1.6.2-bin-hadoop2.6$ jps -l
10355 sun.tools.jps.Jps
9946 org.apache.spark.deploy.master.Master
vagrant@spark-m1:~/spark-1.6.2-bin-hadoop2.6$ tail -n 100
logs/spark-vagrant-org.apache.spark.deploy.master.Master-1-spark-m1.out
```

# Spark WebUI

- `http://spark-m1:8080` master url



## Spark Master at spark://spark-m1:7077

**URL:** spark://spark-m1:7077  
**REST URL:** spark://spark-m1:6066 (*cluster mode*)  
**Alive Workers:** 0  
**Cores in use:** 0 Total, 0 Used  
**Memory in use:** 0.0 B Total, 0.0 B Used  
**Applications:** 0 Running, 0 Completed  
**Drivers:** 0 Running, 0 Completed  
**Status:** ALIVE

### Workers

Worker Id	Address	State	Cores	Memory
-----------	---------	-------	-------	--------

### Running Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

### Completed Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

# Spark Installation - Slave(Worker)

- 共安裝兩台 worker(`spark-s1`, `spark-s2`), 安裝方式相同
- ssh 登入 `spark-s1`
- 安裝 JDK 1.8.x
  - 設定環境變數 `export JAVA_HOME`
  - 將 `$JAVA_HOME/bin` 加入 `PATH` 環境變數
- 修改 `/etc/hosts`, 加入以下資訊

```
172.16.1.10    spark-m1
172.16.1.20    spark-s1
172.16.1.30    spark-s2
```

- 設定 ssh 免密碼登入(`spark-m1 --> spark-s1`)

```
vagrant@spark-s1:~$ cat /vagrant/.ssh/id_rsa.master.pub >> ~/.ssh/authorized_keys
vagrant@spark-s1:~$ sudo su -
root@spark-s1:~$ cat /vagrant/.ssh/id_rsa.master.pub >> ~/.ssh/authorized_keys
```

- 下載 & 解壓縮 Spark 1.6.2

```
vagrant@spark-s1:~$ wget http://d3kbcqa49mib13.cloudfront.net/spark-1.6.2-bin-hadoop2.6.tgz
vagrant@spark-s1:~$ tar zxvf spark-1.6.2-bin-hadoop2.6.tgz
```

# Spark Installation - Slave(Worker)

- 修改 ~/.profile 加入環境變數

```
export SPARK_HOME=/home/vagrant/spark-1.6.2-bin-hadoop2.6
```

- 登出 ssh, 再次登入 spark-s1. 並執行 Spark worker

```
vagrant@spark-s1:~$ cd spark-1.6.2-bin-hadoop2.6
```

```
vagrant@spark-s1:~/spark-1.6.2-bin-hadoop2.6$ sbin/start-slave.sh spark://spark-m1:7077
```

- 檢查是否正常啟動(process & log)

```
vagrant@spark-s1:~/spark-1.6.2-bin-hadoop2.6$ jps -l
```

```
9111 sun.tools.jps.Jps
```

```
9055 org.apache.spark.deploy.worker.Worker
```

```
vagrant@spark-s1:~/spark-1.6.2-bin-hadoop2.6$ tail -n 100
```

```
logs/spark-vagrant-org.apache.spark.deploy.worker.Worker-1-spark-s1.out
```

- 停止 Spark worker

```
vagrant@spark-s1:~/spark-1.6.2-bin-hadoop2.6$ sbin/stop-slave.sh
```

# Spark cluster startup & configuration

- 啟動 cluster `vagrant@spark-m1:~/spark-1.6.2-bin-hadoop2.6$ sbin/start-all.sh`
- 停止 cluster `vagrant@spark-m1:~/spark-1.6.2-bin-hadoop2.6$ sbin/stop-all.sh`
- 重要設定檔:
- `$SPARK_HOME/conf/slaves`
  - 設定所有 worker 的 hostname 資訊
  - 若找不到 slaves 則 Spark 會在 master 另外起一個 worker, 不須另外建立 VM, 方便測試使用.

```
spark-s1
spark-s2
```

- `$SPARK_HOME/conf/spark-defaults.conf`
  - `cp conf/spark-defaults.conf.template conf/spark-defaults.conf`
  - Spark 預設讀取的 conf 檔案

```
spark.master                spark://spark-m1:7077
spark.eventLog.enabled      true
```

- `spark.eventLog.enabled=true` 時, 須要在 master 建立資料夾: `/tmp/spark-events`

<https://spark.apache.org/docs/1.6.2/spark-standalone.html#starting-a-cluster-manually>

# Spark Web UI

- 預設情形下，spark 各項執行資訊只能在執行時期看到
- 保存 application 執行記錄：spark.eventLog.enabled 設為 true



## Spark Master at spark://spark-m1:7077

**URL:** spark://spark-m1:7077

**REST URL:** spark://spark-m1:6066 (*cluster mode*)

**Alive Workers:** 2

**Cores in use:** 2 Total, 0 Used

**Memory in use:** 5.7 GB Total, 0.0 B Used

**Applications:** 0 Running, 7 Completed

**Drivers:** 0 Running, 0 Completed

**Status:** ALIVE

### Workers

Worker Id	Address	State
<a href="#">worker-20160831120609-192.168.1.71-37204</a>	192.168.1.71:37204	ALIVE
<a href="#">worker-20160831120609-192.168.1.72-58342</a>	192.168.1.72:58342	ALIVE



## Execute spark shell

- bin/spark-shell 提供 Scala 的交談式執行環境
- bin/pyspark 提供 Python 的交談式執行環境
- spark-shell, pyspark 內部均呼叫 bin/spark-submit

```
$ bin/spark-shell  
16/09/06 23:52:27 WARN NativeCodeLoader: Unable to load native-hadoop library  
for your platform... using builtin-java classes where applicable  
Welcome to
```



```
version 1.6.2  
  
Using Scala version 2.10.5 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_102)  
Type in expressions to have them evaluated.  
Type :help for more information.  
Spark context available as sc.  
  
SQL context available as sqlContext.  
  
scala>
```

# Execute spark shell - Word count example

```
scala> :paste
// Entering paste mode (ctrl-D to finish)

val words = Seq("copyright", "permission", "software")
val lines = sc.textFile("/vagrant/data/licenses/*.txt")
val counts = lines.flatMap(line=>line.split(" "))
                  .filter( w=>words.contains(w.toLowerCase) )
                  .map( (_, 1) )
                  .reduceByKey( _ + _ )
val result = counts.collect()

// Exiting paste mode, now interpreting.

words: Seq[String] = List(copyright, permission, software)
lines: org.apache.spark.rdd.RDD[String] = /vagrant/data/licenses/*.txt
MapPartitionsRDD[1] at textFile at <console>:29
counts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[5] at reduceByKey
at <console>:34
result: Array[(String, Int)] = Array((software,42), (Software,32),
(SOFTWARE,49), (copyright,67), (Copyright,40), (Permission,15), (permission,15),
(COPYRIGHT,47))

scala> result.foreach(println)
```

# Execute spark example - run-example

- spark 預先打包好的 example jar
  - lib/spark-examples-1.6.2-hadoop2.6.0.jar
- 設定環境變數 *MASTER* 指定 master, 若無則以 local mode 執行
- bin/run-example 內部呼叫 bin/spark-submit

```
$ MASTER=spark://spark-m1:7077 bin/run-example SparkPi 50
16/09/07 00:58:52 WARN NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Pi is roughly 3.142252
$
```

<https://spark.apache.org/docs/1.6.2/#running-the-examples-and-shell>

# Spark Glossary

# Spark Glossary

- Application
  - 使用者開發的 Spark 程式統稱，範圍包含 Driver 與各節點的 Executor
- Driver program
  - 在 Application 裡含有 `main()` 及建立 `SparkContext` 的程式
- Cluster manager
  - 在 cluster 上負責取得可用資源的服務
- Worker node
  - 泛指在 cluster 裡可以執行 application 的節點
- Executor
  - 在 worker 上啟動 application 的 `process(JVM)`。每個 application 都有屬於自己的 Executor

# Spark Glossary

- Job
  - RDD action 所觸發的動作即為一個 Job。每個 Job 均包含多個可平行運算的 task
- Stage
  - 一個 Job 裡的所有 task 會被分解成多組的 task。每一組 task 即為一個 stage
- Task
  - 在 Executor 裡實際執行的任務
- Partition
  - 將 RDD 中的 data 切分為不同的分區，通常一個 partition 的資料由一個 task 處理

# RDD concept & overview

# SparkConf & SparkContext

- *SparkConf* 物件
  - 記錄執行 application(driver) 所需要的資訊
  - 會讀取所有名稱以 'spark.' 開頭的 Java 系統資訊
- *SparkContext* 物件
  - 提供 Driver 執行 Spark Job 的進入點
  - 建立 RDDs, accumulators, broadcast variables
  - run jobs

```
val conf = new SparkConf().setAppName(appName)
                        .setMaster(master)
val sc = new SparkContext(conf)
```



# RDD's trait

- RDD - Resilient Distributed Dataset
  - 顧名思義 - "有彈性的" "分散式" "資料集"
  - 一種物件: `RDD[String]`, `RDD[Int]`
- RDD 特性
  - 容錯, Fault tolerance
  - 平行運算, parallel
  - 不可變, `Immutable(read only)`
  - 資料分區, Partitioned data
  - 記憶體儲存, In-Memory
  - 有資料型別, Typed
  - 快取, Cacheable

# RDD Creation

- 從既有的 collection 建立 RDD: *parallelize*, *makeRDD*

```
scala> val data = Seq(1, 2, 3, 4, 5)  
scala> val distData = sc.parallelize(data)
```

- 由外部讀取: *textFile*
  - Local file system - "file://"


```
scala> val distFile = sc.textFile("data.txt")
```

- HDFS - "hdfs://"
- Amazon S3 - "s3n://"
- Cassandra, HBase

# RDD Operations


- Transformation

- 由現有的 RDD 建立一個新的 RDD
- Lazy, 不立即進行運算
- *map*, *flatMap*, *filter*...



- Action

- 在 RDD 上進行運算, 傳回實際值到 Driver
- *collect*, *count*, *reduce*



String, Int,  
Collection...

# Word count example(again)

```
val words = Seq("copyright", "permission", "software")

val lines = sc.textFile("/vagrant/data/licenses/*.txt")

val counts = lines.flatMap( line=>line.split(" ") )
                    .filter( w=>words.contains(w.toLowerCase) )
                    .map( (_, 1) )
                    .reduceByKey( _ + _ )

val result = counts.collect()
```

# Understand Scala - very, very quickly

```
val words = Seq("copyright", "permission", "software")
```

↳ **val** 宣告一個不可改變的變數, **Seq** 是一個 **scala** 的 **collection** 的物件

```
val lines = sc.textFile("/vagrant/data/licenses/*.txt")
```

↳ **sc** 是 **SparkContext** 物件

```
val counts = lines.flatMap( line=>line.split(" ") )
```

↳ **flatMap** 傳入一個匿名函式, **line** 是傳入值, **split** 傳出 **Array**

```
.filter( w=>words.contains(w.toLowerCase) )
```

↳ **filter** 傳入一個匿名函式, **w** 是傳入值, **contains** 傳出 **Boolean**

```
.map( (_, 1) )
```

↳ **\_** 表示第一個傳入的參數, **(\_,1)** 表示一個 **Tuple2** 物件


```
.reduceByKey( _ + _ )
```

↳ 第一個 **\_** 表示第一個參數, 第二個 **\_** 表示第二個參數

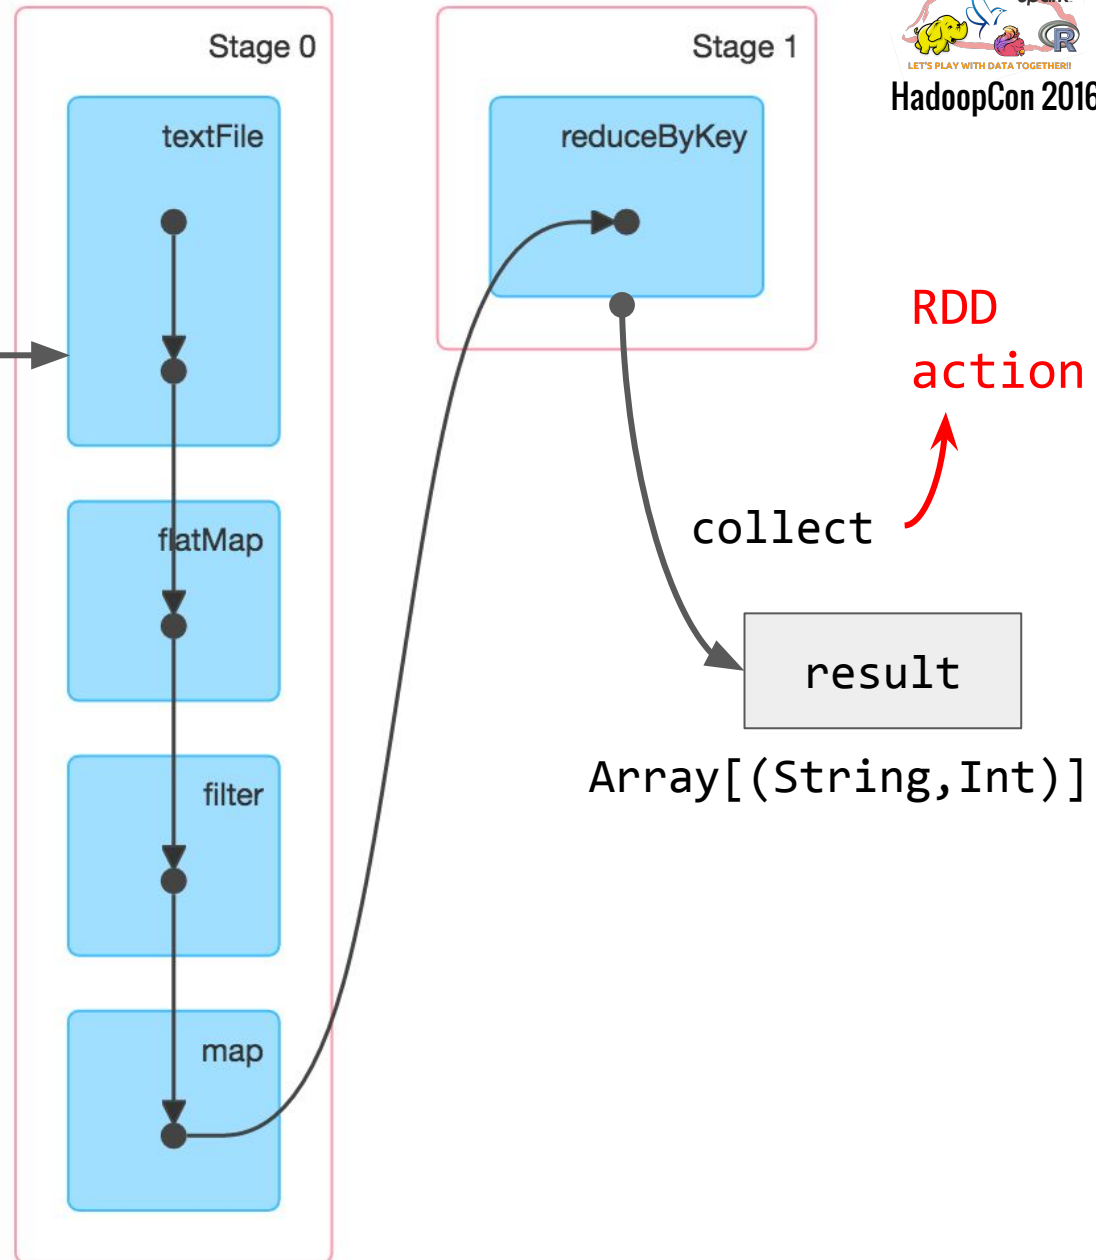
```
val result = counts.collect()
```

↳ **RDD action**

# 觀察 RDD 的變化



DAG Virtualization



# 觀察 RDD 的變化 - Spark shell

RDD.*toDebugString*

```
scala> counts.toDebugString
res1: String =
(35) ShuffledRDD[5] at reduceByKey at <console>:34 []
+- (35) MapPartitionsRDD[4] at map at <console>:33 []
    |   MapPartitionsRDD[3] at filter at <console>:32 []
    |   MapPartitionsRDD[2] at flatMap at <console>:31 []
    |   /vagrant/data/licenses/*.txt MapPartitionsRDD[1] at textFile at <console>:29 []
    |   /vagrant/data/licenses/*.txt HadoopRDD[0] at textFile at <console>:29 []

scala>
```

# RDD Persistence

- RDD.*persist()* , RDD.*cache()*
- RDD 的結果保留在記憶體內，下次使用不須重新計算

```
val words = Seq("copyright", "permission", "software")
```

```
val lines = sc.textFile("/vagrant/data/licenses/*.txt")
```

```
val allWords = lines.flatMap(line=>line.split(" "))
```

```
allWords.persist()
```

```
val counts = allWords.filter( myfilter1 )  
                      .map( (_, 1) ).reduceByKey( _ + _ )
```

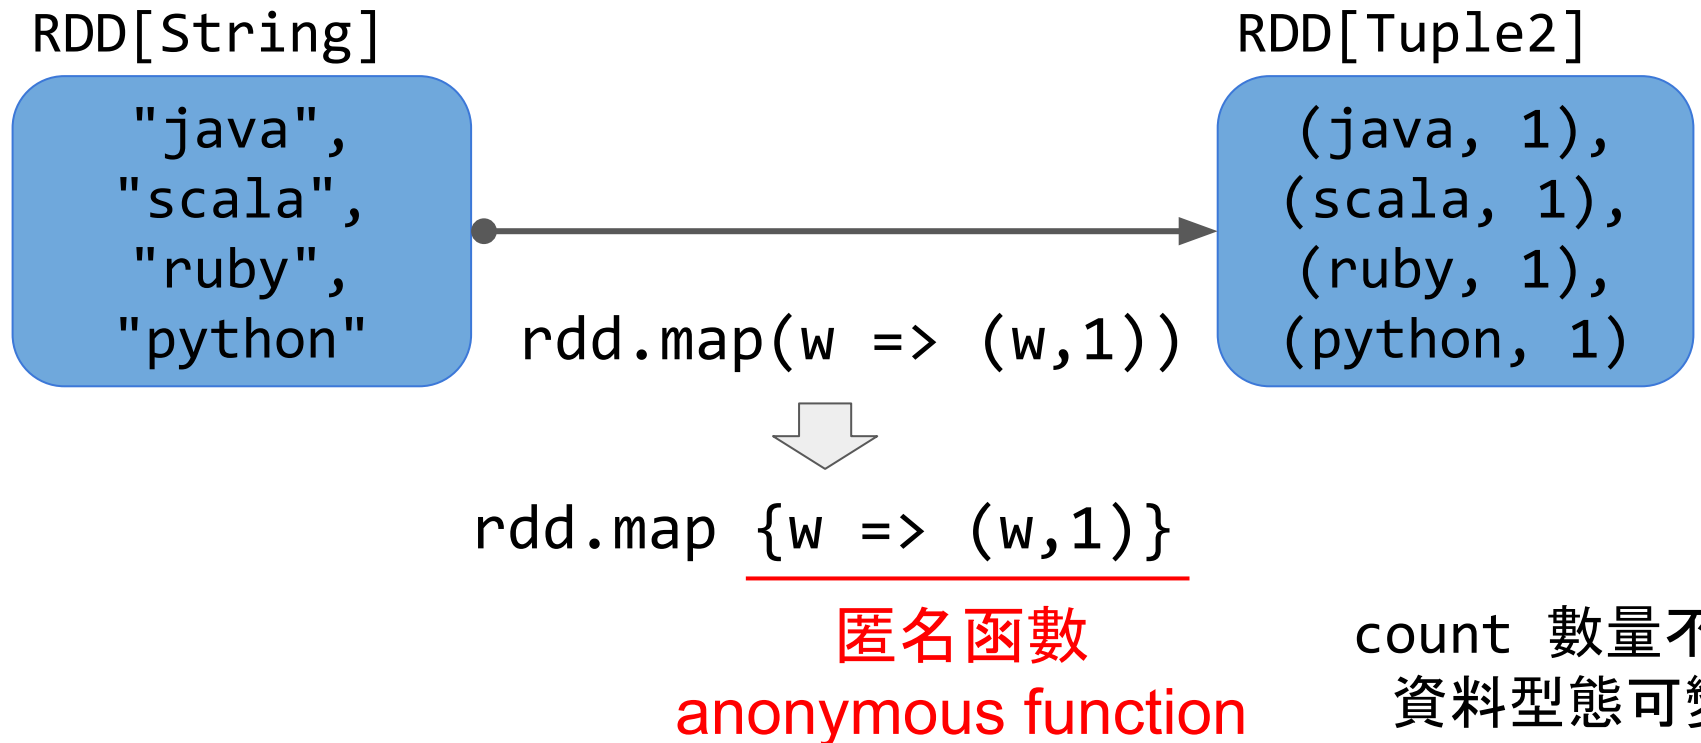
```
val counts2 = allWords.filter( myfilter2 )  
                      .map( (_, 1) ).reduceByKey( _ + _ )
```



# RDD operations - transformation

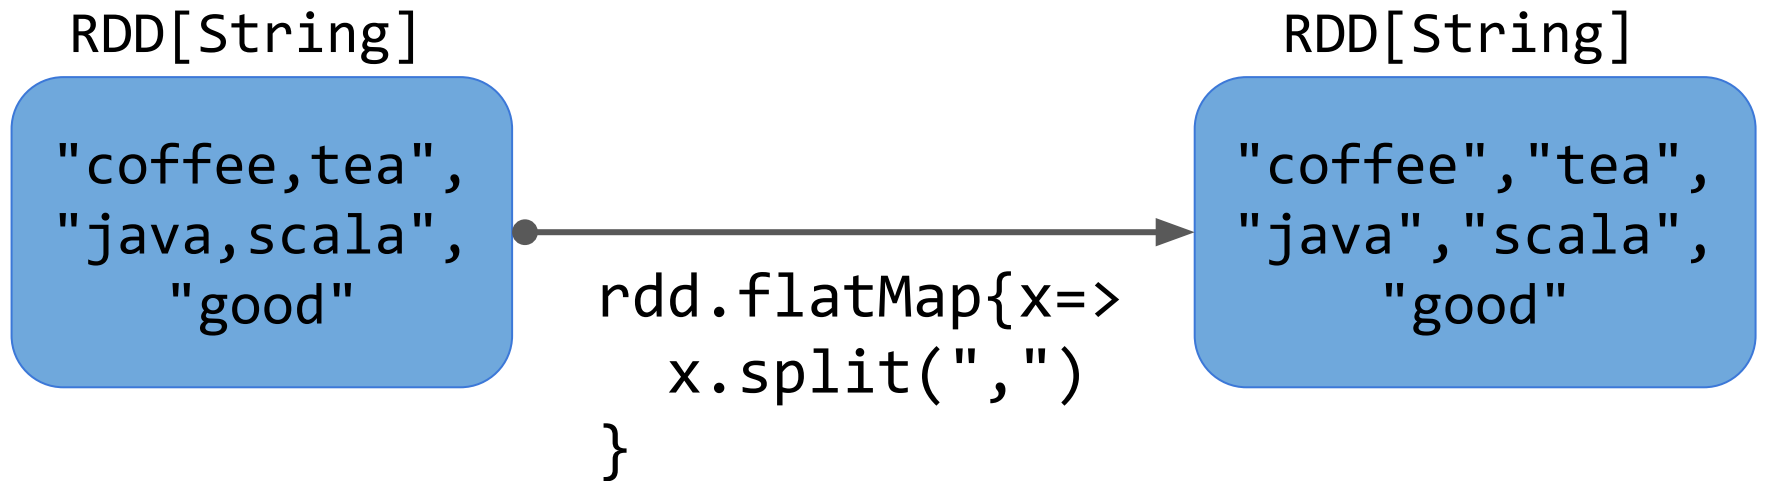
# RDD.map()

Return a new RDD by applying a function to all elements of this RDD.



# RDD.flatMap()

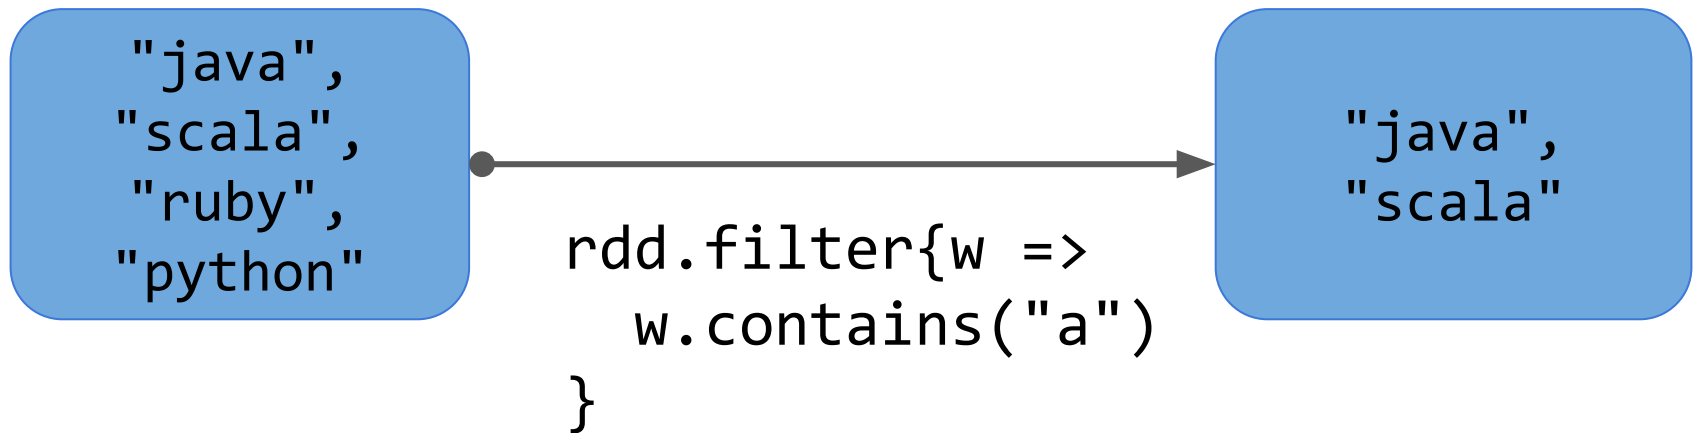
Return a new RDD by first applying a function to all elements of this RDD, and then flattening the results.



count 數量可變  
資料型態可變

# RDD.filter()

Return a new RDD containing only the elements that satisfy a predicate.



count 數量改變  
資料型態不變

# Key-Value Pairs

- Scala 物件: *Tuple2*
  - Scala 三種不同寫法
  - 取第一個值: \_1
  - 取第二個值: \_2

```
scala> val t1 = ("java", 100)
t1: (String, Int) = (java,100)

scala> val t2 = "java" -> 100
t2: (String, Int) = (java,100)

scala> val t3 = new Tuple2("java", 100)
t3: (String, Int) = (java,100)

scala> println(t1._1 + "/" + t1._2)
java/100
```

- *Tuple2* 裡的 \_1 作為 key, \_2 作為 value
- Key-Value Pairs 最常見的應用場景是在進行 group, aggregation 等需要 shuffle 的操作, 如 *groupByKey*, *reduceByKey*, *join* ...等
- 參考 API: [PairRDDFunctions](#)

# RDD.reduceByKey()

Merge the values for each key using an associative reduce function. This will also perform the merging locally on each mapper before sending results to a reducer, similarly to a "combiner" in MapReduce. Output will be hash-partitioned with the existing partitioner/ parallelism level.

RDD[(String, Int)]

```
(java, 1),
(scala, 1),
(python, 1),
(scala, 1)
```

RDD[(String, Int)]

```
(java, 1),
(scala, 2),
(python, 1)
```

```
rdd.reduceByKey{(v1,v2)=>
  v1 + v2
}
```

=

```
rdd.reduceByKey(_ + _)
```

v1

v2

# RDD operations - action

# RDD actions

Action	Meaning
<b>reduce(<i>func</i>)</b>	Aggregate the elements of the dataset using a function <i>func</i> (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.
<b>collect()</b>	Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.
<b>count()</b>	Return the number of elements in the dataset.
<b>first()</b>	Return the first element of the dataset (similar to take(1)).



# RDD actions

Action	Meaning
<b>saveAsTextFile(<i>path</i>)</b>	Write the elements of the dataset as a text file (or set of text files) in a given directory in the local filesystem, HDFS or any other Hadoop-supported file system. Spark will call <code>toString</code> on each element to convert it to a line of text in the file.
<b>countByKey()</b>	Only available on RDDs of type (K, V). Returns a hashmap of (K, Int) pairs with the count of each key.
<b>foreach(<i>func</i>)</b>	Run a function <i>func</i> on each element of the dataset. This is usually done for side effects such as updating an <a href="#">Accumulator</a> or interacting with external storage systems.

# More transformations and actions

- Spark programming guide - [Transformations](#)
- Spark programming guide - [Actions](#)
- Unit Test - 最佳學習材料
  - [org.apache.spark.rdd](http://org.apache.spark.rdd)


# Shuffle & RDD Dependency

# Shuffle

- 資料重組 — Shuffle 是一種重新分配 data 的機制，跨越不同的 partition 進行資料重組
- 通常會涉及跨越 Executor 或實體機器的資料複製
- 代價高昂
  - 磁碟 I/O
  - 資料序列化
  - 網路 I/O
- Shuffle 通常是 spark job performance tuning 的重點

# Shuffle


ReduceByKey



[https://databricks.gitbooks.io/databricks-spark-knowledge-base/content/best\\_practices/prefer\\_reducebykey\\_over\\_groupbykey.html](https://databricks.gitbooks.io/databricks-spark-knowledge-base/content/best_practices/prefer_reducebykey_over_groupbykey.html)

# RDD Dependency

- Narrow dependency
  - child RDD 只依賴 parent RDD 固定的 partition
- Wide dependency
  - child RDD 每一個 partition 均依賴 parent RDD 所有的 partition



# Thank you !